BFCL Benchmark — Model Comparison (Apple M4 Pro, Metal)

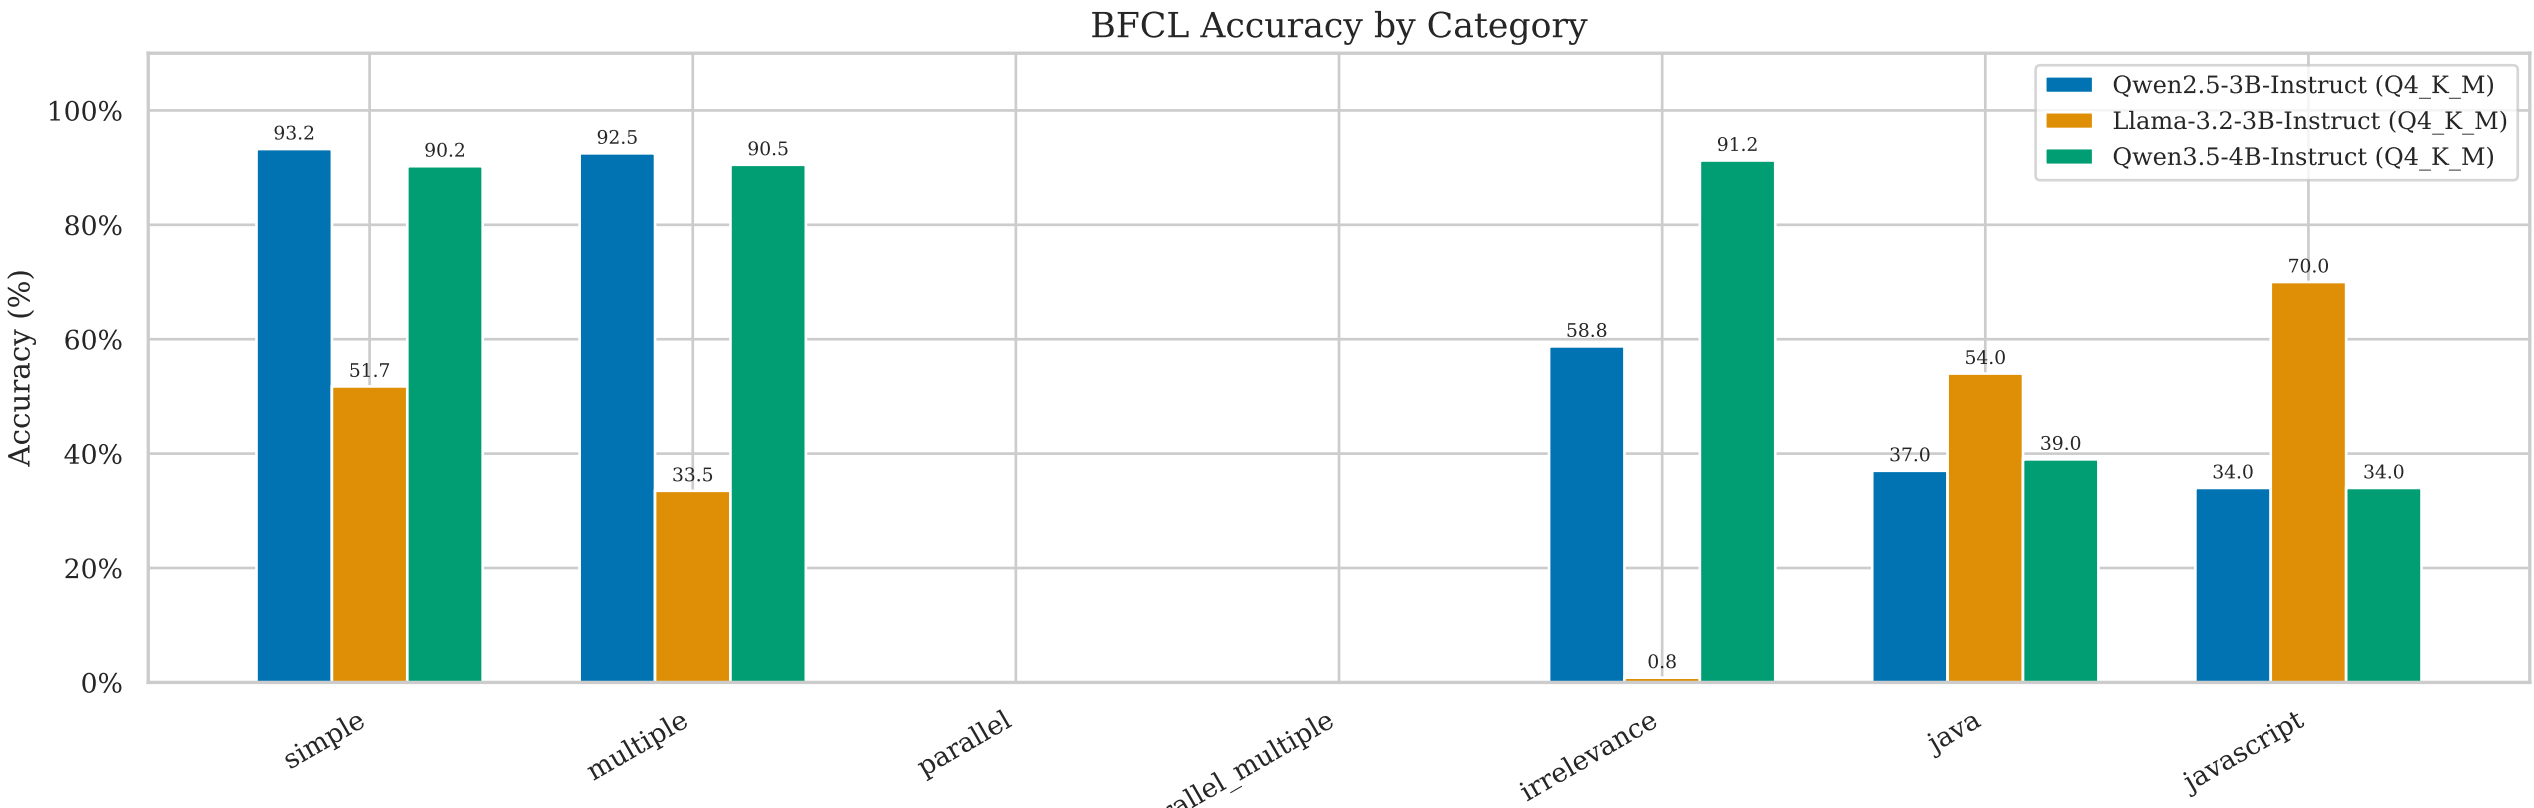

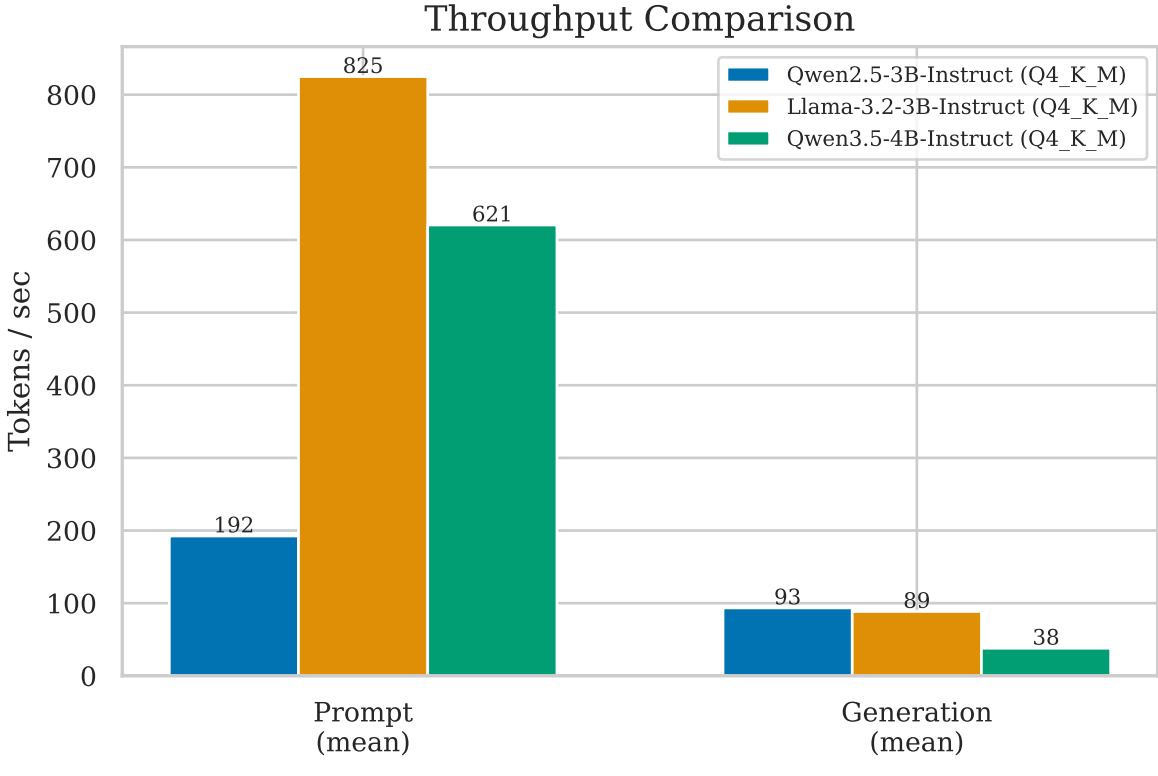

Summary

Model	Avg Acc	Prompt tok/s	Gen tok/s	TTFT (ms)	Peak RSS (MB)
Qwen2.5-3B-Instruct (	45.1%	192	93	12	8547
Llama-3.2-3B-Instruct (	30.0%	825	89	218	12384
Qwen3.5-4B-Instruct (	49.3%	621	38	656	4075