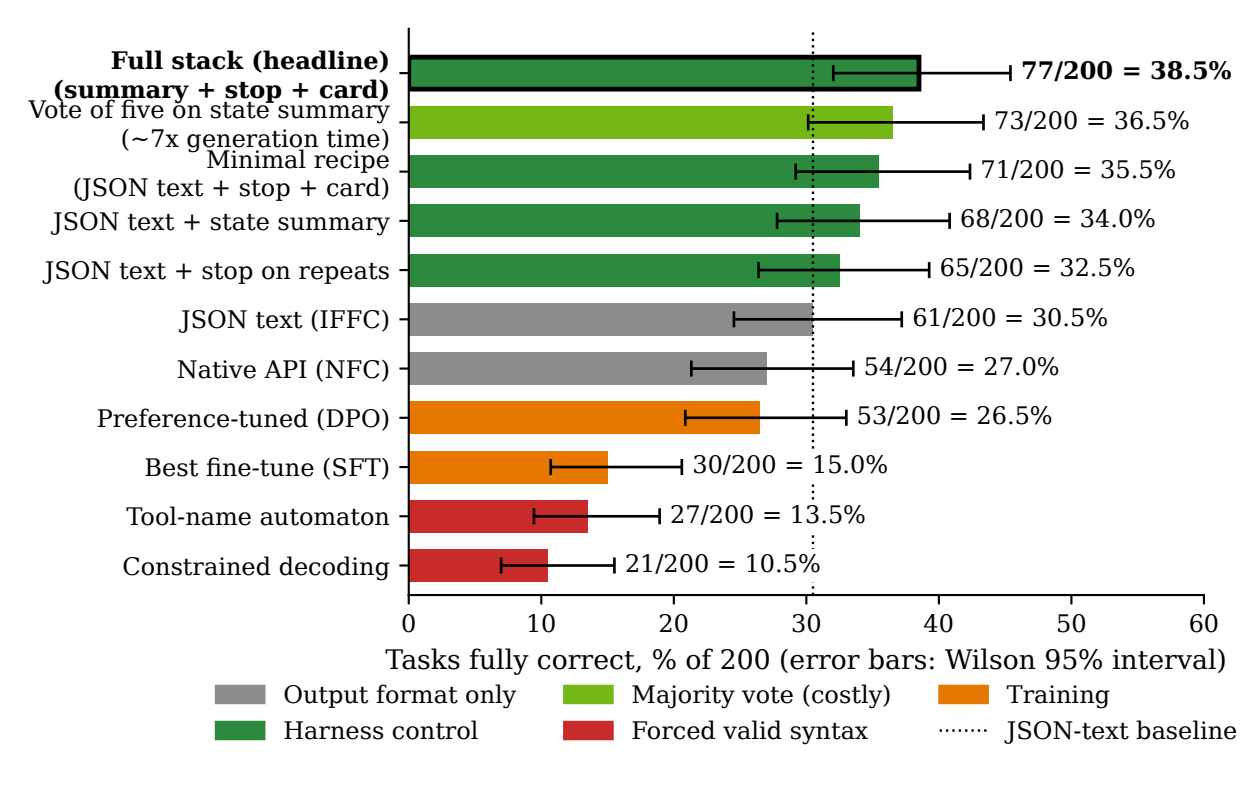
Full stack (headline)
(summary + stop + card)
77/200 = 38.5%
Vote of five on state summary
(~7x generation time)
73/200 = 36.5%
Minimal recipe
(JSON text + stop + card)
71/200 = 35.5%
JSON text + state summary
68/200 = 34.0%
JSON text + stop on repeats
65/200 = 32.5%
JSON text (IFFC)
61/200 = 30.5%
Native API (NFC)
54/200 = 27.0%
Preference-tuned (DPO)
53/200 = 26.5%
Best fine-tune (SFT)
30/200 = 15.0%
Tool-name automaton
27/200 = 13.5%
Constrained decoding
21/200 = 10.5%
0
10
20
30
40
50
60
Tasks fully correct, % of 200 (error bars: Wilson 95% interval)
Output format only
Majority vote (costly)
Training
Harness control
Forced valid syntax
JSON-text baseline