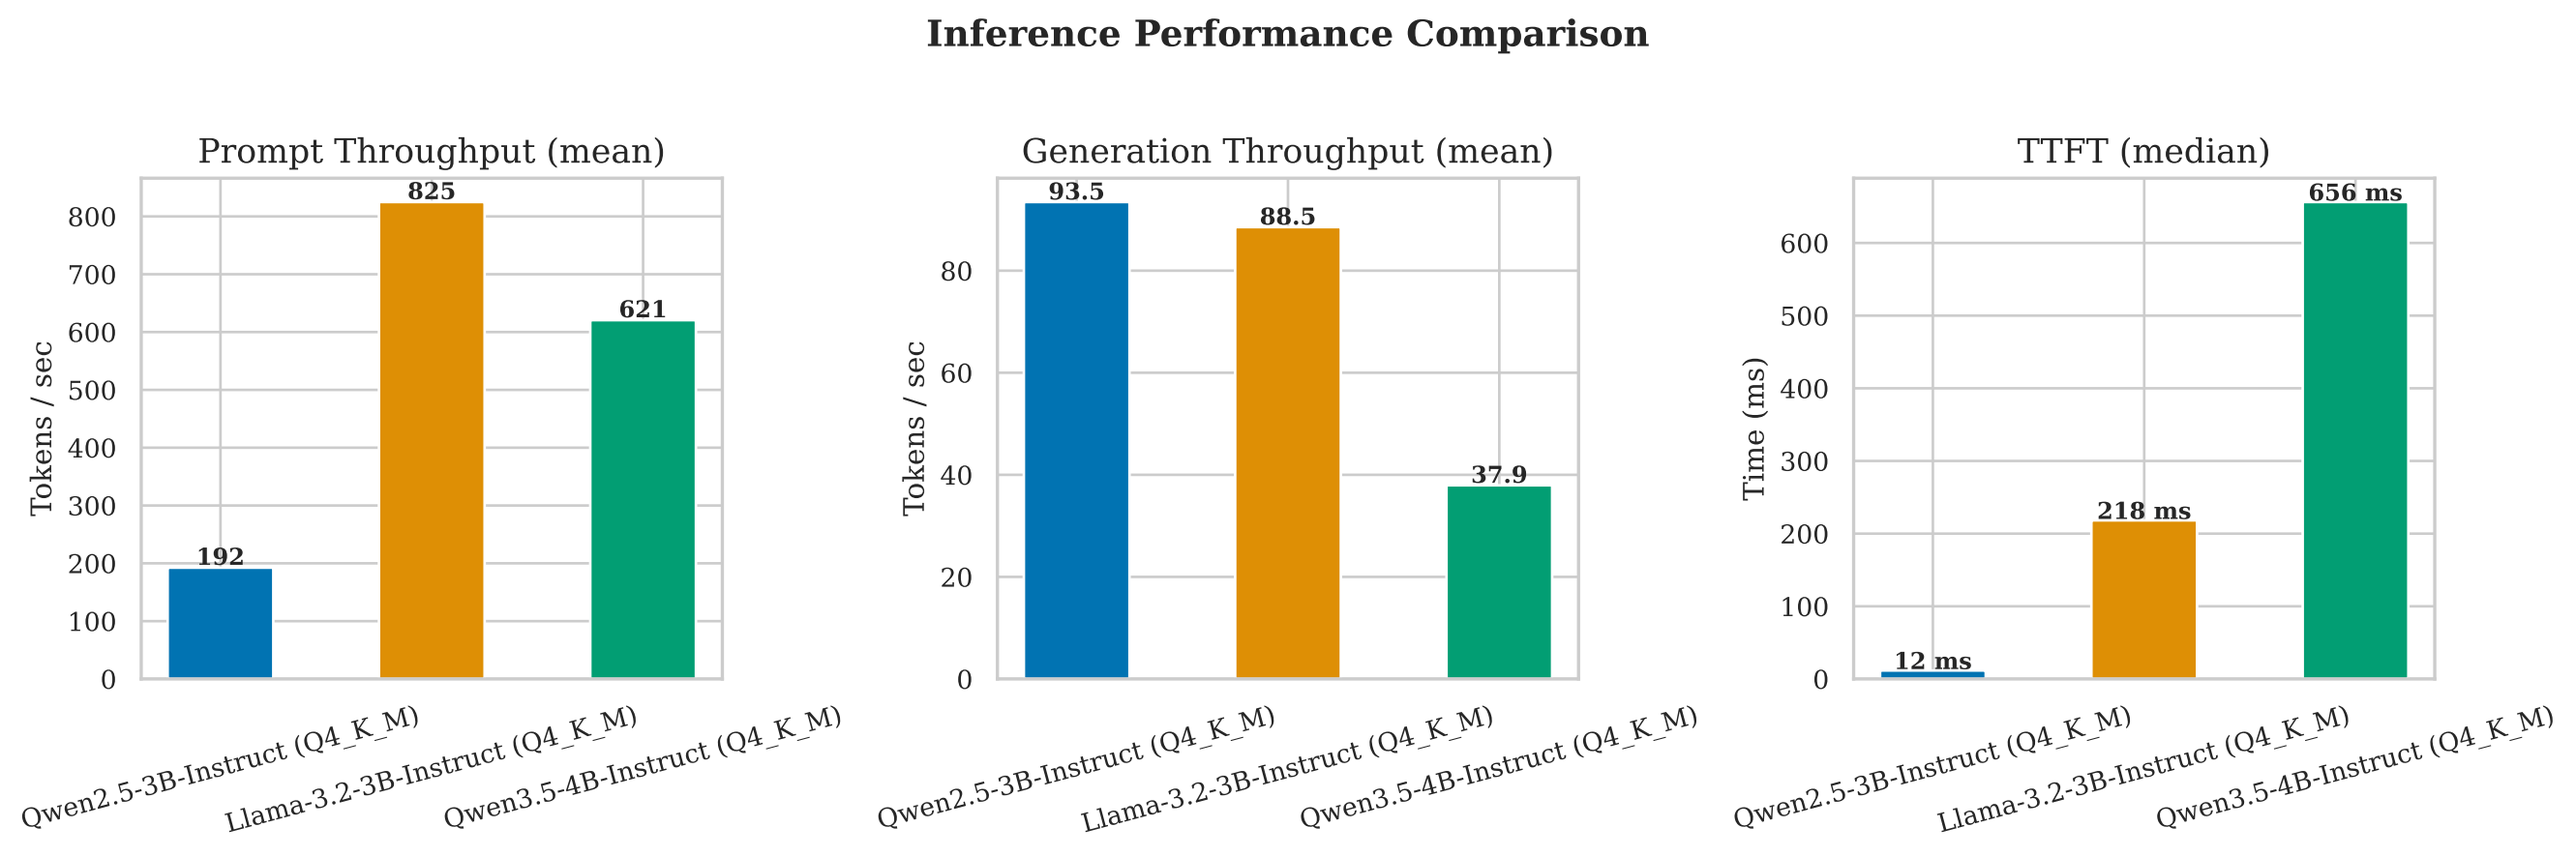
Inference Performance Comparison
Prompt Throughput (mean)
Tokens / sec
0
100
200
300
400
500
600
700
800
192
825
621
Qwen2.5-3B-Instruct (Q4_K_M)
Llama-3.2-3B-Instruct (Q4_K_M)
Qwen3.5-4B-Instruct (Q4_K_M)
Generation Throughput (mean)
Tokens / sec
0
20
40
60
80
93.5
88.5
37.9
Qwen2.5-3B-Instruct (Q4_K_M)
Llama-3.2-3B-Instruct (Q4_K_M)
Qwen3.5-4B-Instruct (Q4_K_M)
TTFT (median)
Time (ms)
0
100
200
300
400
500
600
12 ms
218 ms
656 ms
Qwen2.5-3B-Instruct (Q4_K_M)
Llama-3.2-3B-Instruct (Q4_K_M)
Qwen3.5-4B-Instruct (Q4_K_M)