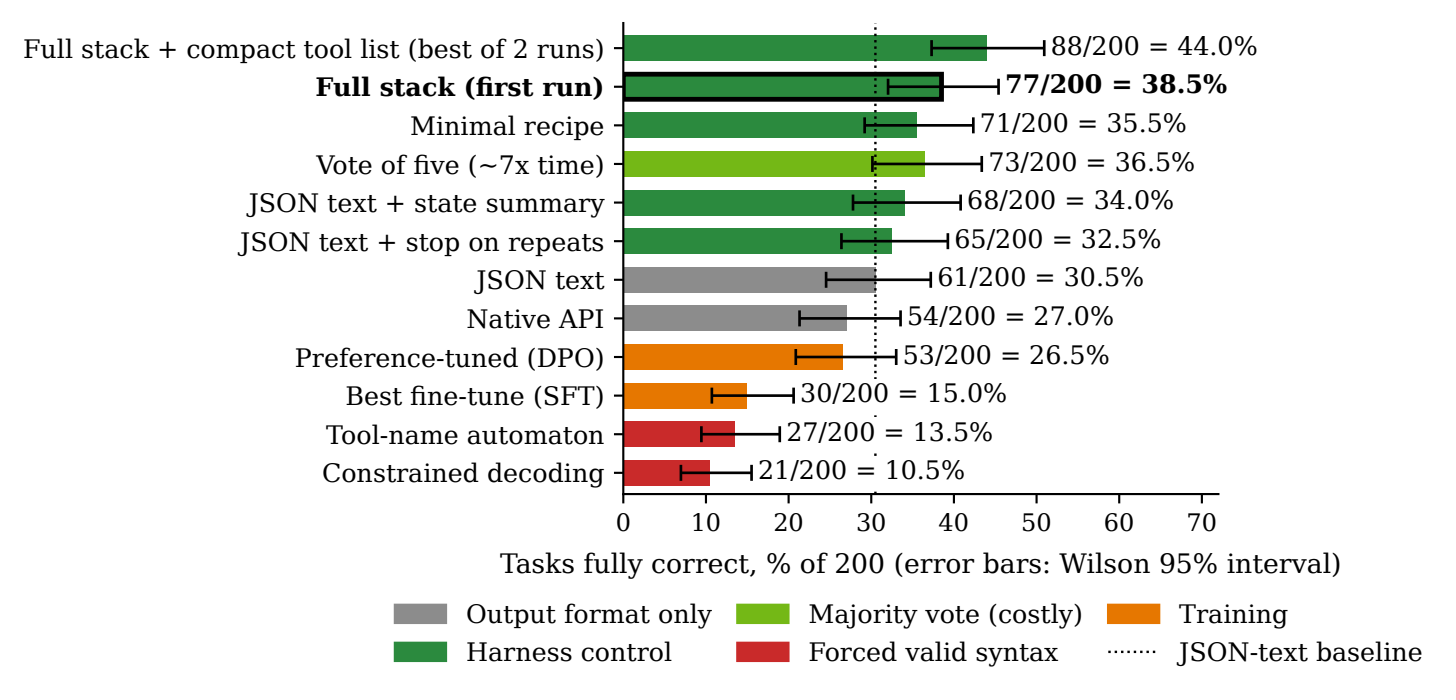
Full stack + compact tool list (best of 2 runs)
Full stack (first run)
Minimal recipe
Vote of five (~7x time)
JSON text + state summary
JSON text + stop on repeats
JSON text
Native API
Preference-tuned (DPO)
Best fine-tune (SFT)
Tool-name automaton
Constrained decoding
88/200 = 44.0%
77/200 = 38.5%
71/200 = 35.5%
73/200 = 36.5%
68/200 = 34.0%
65/200 = 32.5%
61/200 = 30.5%
54/200 = 27.0%
53/200 = 26.5%
30/200 = 15.0%
27/200 = 13.5%
21/200 = 10.5%
0
10
20
30
40
50
60
70
Tasks fully correct, % of 200 (error bars: Wilson 95% interval)
Output format only
Harness control
Majority vote (costly)
Forced valid syntax
Training
JSON-text baseline